

The SaaS Apocalypse Fallacy:

Why AI Raises the Value of the SaaS Bundle

Yakov Pyotr Shkolnikov
Independent Researcher

April 20, 2026

Executive Summary. In the first week of February 2026, more than one trillion dollars in market capitalization was erased from software stocks. The sell-off carried a specific claim: AI will replace enterprise Software-as-a-Service, and the economics of enterprise software procurement will invert. This paper examines the procurement economics of that claim. It argues that the SaaSocalypse narrative rests on a category error. AI in its current and near-term delivery form is not an alternative to SaaS. It is SaaS, delivered under contracts that are structurally weaker on identifiable dimensions. Enterprises that replace mature SaaS with in-house AI builds absorb five structural asymmetries the replacement calculus does not count: liability, compliance generation, cost concentration, expertise and maintenance, and integration and data responsibility. For the 2026–2030 enterprise procurement horizon, AI raises the value of the SaaS bundle because the deployer now sees what the bundle was quietly covering.

1. The Category Error

In the first week of February 2026, software stocks lost more than \$1 trillion in market capitalization [1]. Don Muir reported that approximately \$300 billion of market value evaporated in a single trading day on February 4, with named enterprise-software incumbents including Salesforce, ServiceNow, Adobe, and Workday each dropping around 7 percent and Intuit falling nearly 11 percent; the IGV Software Index had already declined roughly 30 percent from its late-September peak [2]. Andreessen Horowitz used the term “SaaSocalypse” while arguing in the same essay that AI will expand rather than destroy the software industry, locating software value in durable moats rather than in code [3]. The market narrative driving the sell-off was blunter: AI will replace enterprise SaaS, and the economics of enterprise software procurement will invert.

This paper assumes AI is being deployed and examines the procurement economics of that deployment. It is not a critique of AI capability or of AI integration by SaaS vendors. Three terms recur with stable meanings: *SaaS* (a mature multi-tenant service under commercial contract), *AI-as-SaaS*

(a vendor-delivered AI service consumed through an API, agent platform, or coding tool), and *in-house AI build* (an application the enterprise builds itself, typically on an AI-as-SaaS dependency).

The SaaSocalypse narrative rests on a category error: evaluating AI against the capability line of a SaaS product while ignoring the rest of the bundle that product was delivering. In its current and near-term delivery form, AI is SaaS, delivered by vendors, priced on consumption, and governed by contracts that are structurally weaker than the SaaS contracts enterprises already hold. An enterprise that replaces a mature HR SaaS with an in-house AI build has not exited a vendor relationship. It has entered a different one while absorbing five structural asymmetries that the SaaS relationship was quietly covering.

An enterprise pursuing this replacement has four paths. None escapes AI-as-SaaS dependency.

The first path is direct AI API consumption. An application that calls a frontier model API is a consumption-priced AI-as-SaaS with weaker contracts than the SaaS it replaced.

The second path is an in-house AI build written with AI coding tools. The code may not call AI APIs at runtime, but the organization that wrote it with AI assistance has a maintenance capability defined by continued access to those tools. The AI dependency moved from the runtime to the maintenance loop, and the coding tools themselves are AI-as-SaaS.

The third path is self-hosted open-weight deployment. Open-weight frontier models reached parity with their closed counterparts across multiple independent model families in 2025 and 2026. Operating these models at enterprise throughput requires hyperscaler-scale infrastructure. For most enterprises, the practical deployment pattern is managed hosting through a major cloud provider's AI platform or a specialist open-weight host. That pattern returns the enterprise to Path 1 with deployer-selected models. Open-weight models delivered inside a SaaS wrapper restore every bundle component the paper will describe: support, liability, compliance posture, amortized infrastructure, maintained expertise. For most enterprises, the choice between open-weight and closed-model access is rarely a clean exit from SaaS.

The fourth path does not require the enterprise to build at all: AI delivered as outcome-based services. Julien Bek's March 2026 Sequoia essay "Services: The New Software" describes this as "an autopilot sells the work" [4]. That remains a vendor relationship under contract. The five asymmetries apply at varying intensities, with the cost asymmetry reshaped by the services vendor's variable performance-risk pricing.

This paper addresses the 2026–2030 procurement horizon for enterprises that consume software. It is not written for entities that build or distribute AI as their product — SaaS vendors integrating AI into their offerings, AI tool developers, frontier model labs, or hyperscalers operating at infrastructure scale. Those organizations have different procurement economics because AI is what they build, not what they consume. The paper also does not address deployers operating under sovereignty-mandated frameworks such as EU data residency, defense, or sector-specific sovereign-AI contexts. Those environments are out of scope.

This is a build-versus-buy decision, and several of the asymmetries below apply to any custom build. What is AI-specific is that AI coding tools make the build path look dramatically cheaper and

faster than it is, which encourages organizations to skip the discipline a normal build-versus-buy analysis would impose. Two of the asymmetries (liability from unbounded probabilistic output, and maintenance dependency on AI tooling that is itself SaaS) are meaningfully AI-specific. The others sharpen with AI rather than originate with it.

Consider a CIO weighing a proposal to replace the enterprise HR SaaS with an in-house AI build. The proposal estimates the build at 18 months with a team of two to five engineers, AI-assisted throughout. It projects per-seat savings against the SaaS license and operational savings from exiting the vendor relationship. These are the visible line items: the same ones the SaaS apocalypse narrative counts. The rest of this paper examines what the proposal and the narrative leave uncounted.

2. The SaaS Bundle

SaaS, though often treated as a capability, is better operationalized as a bundle of components that enterprises rarely itemize because the vendor delivers them automatically. The structure of the bundle is a transaction-cost-economics argument (after Williamson [5]): enterprise procurement of specialized software involves asset specificity, vendor opportunism, and bounded rationality about integration cost, and the SaaS bundle is a governance response to those hazards. Figure 1 visualizes this disaggregation.




Figure 1: **The SaaS bundle disaggregated.** SaaS vendors absorb six components of the bundle on the enterprise’s behalf. AI-as-SaaS vendors sell capability and leave the other five components to the deployer; vendor-absorbed components are solid and blue-filled, deployer-generated components are dashed and red-filled. The SaaS apocalypse narrative evaluates AI against the capability row alone.

The SaaS apocalypse narrative evaluates AI against the capability row and reads an advantage. The comparison is incomplete: the other five rows are where the uncounted structural cost lives,

and for each of them the AI-as-SaaS contract shifts the component from vendor-absorbed to deployer-generated, explicitly, in the contract language. AI has disaggregated the bundle rather than eliminating it, and in doing so has made visible what the bundle was always quietly doing. Each of the five non-capability components is examined below, and the whole set is summarized in Figure 2.

3. Five Asymmetries




Figure 2: **The five asymmetries.** Each procurement dimension shifts from vendor-absorbed under mature SaaS to deployer-generated under AI-as-SaaS, with the primary evidence anchor noted beneath.

3.1 Asymmetry 1 — Liability

Mature SaaS contracts allocate defined failure modes (uptime, data loss, security breach, limited IP infringement) to the vendor under contractual remedies. AI-as-SaaS contracts allocate to the deployer what this paper refers to as *output-harm risk*: the consequences of an AI-generated output being wrong, misleading, discriminatory, defamatory, infringing, or otherwise harmful when used by the deployer’s customers, employees, or counterparties. The deployer is treated in these contracts as the responsible publisher of the model’s output, regardless of how the output was produced.

SaaS baseline protections include conformance warranties where negotiated, IP indemnification at enterprise-tier coverage rates, and audit-rights and DPA provisions in typical commercial agreements. Output-quality warranties for AI systems are routinely absent. Liability caps at fees paid are comparable between SaaS and AI in the aggregate, which masks material differences on specific contract dimensions. The *TermScout 2026 Contract Trust Report*, summarized in the evidence box below, documents these differences at the clause level.

TermScout's *2026 Contract Trust Report* is based on clause-level analysis of thousands of 2025 commercial agreements. It documents a measurable compression of AI-related indemnification scope between 2022 and 2025, with contracts increasingly breaking unified IP indemnities into fragmented structures across input, training, and output rights. Insurance retreated from AI risk during 2025 as hallucinations, model drift, and training-data contamination proved difficult to price; risk management shifted into the contract itself through narrower indemnities and performance-linked liability. The report observes systematic divergence between vendor-drafted and buyer-drafted AI provisions, with vendors framing compliance obligations narrowly and buyers pushing for broader visibility and event-triggered rights [6].

These are observations about the current state of the AI procurement category, not claims about vendor intent. Some vendors offer above-market indemnification, but these programs are narrow in scope and conditional on mitigations. The category as a whole has an immature risk-allocation profile relative to the SaaS baseline: SaaS contracts are imperfect risk-transfer mechanisms, and AI-as-SaaS contracts at present shift substantially more risk to the deployer than the SaaS contracts they would replace.

3.2 Asymmetry 2 — Compliance Generation

SaaS certifications cover the certified system end-to-end: a SaaS vendor's SOC 2 Type II applies to the HR system the enterprise uses, and the vendor's DPA governs the vendor's processing of employee data inside that system. When internal audit asks whether the HR system is compliant, it points to the vendor's report and inherits a posture. AI API vendor certifications, by contrast, apply to the API infrastructure only; the enterprise application built on top of that API requires its own compliance posture, generated from scratch.

The generation burden applies to every enterprise deployer, regardless of sector: SOC 2, ISO 27001, HIPAA where health data is involved, PCI-DSS where payment data is involved, GDPR where EU data subjects are involved, and the corresponding US sector-specific frameworks all extend to the deployer's AI application layer rather than stopping at the vendor's API. The EU AI Act is the strongest, most explicit formal statement of the pattern and is shown in the anchor below; enterprises operating in less heavily regulated contexts face a lighter but recognizably similar list rather than an absence of obligation.

The EU AI Act, Article 26, enumerates specific deployer obligations for high-risk AI systems: appropriate technical and organizational measures, human oversight assignment, input data representativeness, log retention for at least six months, incident notification to providers and authorities, and notification of workforce where applicable [7]. Deployer obligations at this level of granularity have limited SaaS precedent.

Cloud Security Alliance guidance on AI organizational responsibilities documents the more general security considerations that expand the enterprise audit scope for any deployer, regulated or not: input and output safety (protecting against malicious or harmful inputs and outputs), model

lifecycle and management oversight, continuous monitoring with detection of data drift and model drift, and auditing mechanisms to track AI system activity [8]. Compliance posture now must be generated for the deployer’s application, component by component. The deployer generates the certification; they do not inherit it.

3.3 Asymmetry 3 — Cost Concentration

SaaS pricing amortizes development, compliance, security, support, and maintenance across the vendor’s shared customer base. An enterprise running a single-tenant in-house AI build bears those costs alone, on its own balance sheet, without the amortization benefit of a multi-tenant customer base underneath.

The apparent cost comparison is AI API cost per transaction against SaaS per-seat license, which is the comparison the apocalypse narrative makes. A total-cost-of-ownership comparison adds single-tenant compliance generation, security audit, maintenance, liability reserve, integration layer, and model-change response to the AI API line. This paper does not compute the TCO delta; the delta exists, is systematically undercounted, and AI API pricing does not represent it.

A related dimension is predictability. SaaS per-seat pricing produces forecastable spend. AI procurement pricing is consumption-based (per token, per call, per conversation, per resolution), and frontier vendors and agentic entrants are moving further from per-seat toward usage- and outcome-denominated pricing [3]. Unit economics become knowable only after consumption, and forecasting AI-application spend requires workload modeling that the SaaS license did not require.

Uber’s full-year 2026 AI budget was maxed out within months, driven by rapid adoption of AI coding tools. CTO Praveen Neppalli Naga told *The Information*: “I’m back to the drawing board because the budget I thought I would need is blown away already.” The company maintains internal leaderboards ranking engineer AI use [9].

MIT’s NANDA project, in *The GenAI Divide: State of AI in Business 2025*, found that external partnerships succeed at approximately twice the rate of internal builds, and that only five percent of custom enterprise AI tools reach production [10]. A CIO quoted in the report captured the pattern: “We’ve seen dozens of demos this year. Maybe one or two are genuinely useful. The rest are wrappers or science projects.”

Chuck Fuerst, quoted in Steve Banker’s March 2026 *Forbes* piece, described the cost-surface pattern at the engineering level: “What AI accelerates is the visible layer: the UI, the workflow, the initial functionality. What it doesn’t solve is everything underneath: architecture, reliability, integration, compliance, and long-term maintainability” [11].

Apparent cost and total cost of ownership are therefore different numbers. The gap between them is the sum of the non-capability bundle components that the SaaS license was covering and the AI-as-SaaS contract is not. The replacement decision rarely itemizes it.

3.4 Asymmetry 4 — Expertise and Maintenance

A mature SaaS vendor spreads expertise across a team: engineers, domain specialists, and compliance staff whose full-time work is maintaining the system the enterprise consumes. When individuals leave, the vendor’s institutional processes absorb the departure through onboarding, runbooks, tests, and overlapping ownership. In-house AI builds typically concentrate expertise in a small number of engineers; when they leave, the system continues to run but its behavior becomes less legible, and the organization’s maintenance capability reduces to its ability to direct AI coding tools to modify code those tools originally produced. AI-assisted code generation, without compensating review and documentation discipline, can widen rather than close this concentration. The evidence anchor below reports one early measurement of the comprehension gap such workflows can produce.

Anthropic’s research has begun to quantify the comprehension gap that AI-assisted development can introduce. In a randomized controlled trial of 52 software engineers reported by Shen and Tamkin, developers using AI coding assistance scored 17 percent lower on a subsequent comprehension quiz than those who completed the same task without AI assistance, with no significant improvement in task completion time [12]. The result is suggestive rather than definitive, and it applies most sharply to teams that do not compensate with stronger review, testing, and documentation practices.

Osmani articulated the broader phenomenon as *comprehension debt*: the gap between the code that exists in a system and the code any human in the organization genuinely understands [13]. The paper applies it to the SaaS-replacement decision: when an enterprise replaces a SaaS system by building an AI application with AI coding assistance, its maintenance capability is defined by continued access to AI tooling that is itself AI-as-SaaS, procured under contracts that exhibit the liability asymmetry described above.

This is different in kind from inherited legacy code. With legacy systems the original authors reasoned through the decisions and can often still be consulted. With AI-assisted codebases the reasoning path that produced the code lives only in the tool session and is harder to reconstruct after the fact.

3.5 Asymmetry 5 — Integration and Data Responsibility

Inside a mature SaaS product, the vendor owns integration with surrounding systems and the interpretation of the data that flows through it (field-level semantics, business-rule enforcement, reconciliation logic). AI applications that replace SaaS transfer that ownership to the deployer in two directions.

When data stays in existing systems and an AI layer reads from them, the enterprise owns the entire pipeline between the data and the model: schema management, validation, field-level interpretation, and business-rule enforcement all become application-layer work. When data migrates out of the existing systems into an AI-native application, the interpretation layer the SaaS

vendor had encoded in its schemas and business logic does not migrate with it; re-encoding that layer is substantial and routinely underestimated work. The evidence anchor below reports two complementary findings about how often this burden stalls projects.

Gartner predicts that through 2026, 60 percent of AI projects that lack adequate data readiness will be abandoned [14]. This is a conditional failure rate within the data-unready subset, not a projection across AI projects generally. MuleSoft’s 2025 Connectivity Benchmark Report, based on interviews with 1,050 IT leaders globally, finds that “nearly all (95%) of IT leaders report integration as a hurdle to implementing AI effectively” [15].

Model drift compounds the deployer’s data-and-integration burden. A deployed model’s behavior changes when the vendor updates it. Updates happen on vendor schedules, often without any contractual obligation to notify the deployer. A customer segmentation that produced one output distribution in January may produce a different distribution in March on identical input data. Detecting this drift, responding to it, and reconciling downstream systems with the new behavior all become deployer responsibilities, whereas a mature SaaS vendor manages this inside its own change-management scope.

When AI replaces SaaS but data stays, the deployer owns the new integration stack. When AI replaces SaaS and data migrates, the deployer owns the interpretation layer. Either way, integration and data-interpretation responsibility transfer from the vendor to the deployer.

4. Conclusion

Five independent 2026 voices converge on a common observation about the February 2026 market event. Forrester forecasts traditional SaaS spending growing from \$318 billion in 2025 to \$576 billion in 2029, and concludes that the enterprise core remains while the central nervous system evolves [1]. Muir frames the market drop as a “reallocation of where value accrues inside a much larger market” and cites a Goldman Sachs projection that more than 60 percent of software economics could flow through agentic systems by 2030 [2]. Andreessen Horowitz identifies the moats that remain: network effects, brand, cornered resources, and process power, the accumulated institutional workflow knowledge of mature SaaS vendors [3]. Sequoia’s David Cahn observes that enterprises struggling with in-house AI implementation are accelerating startup momentum rather than diminishing software demand [16]. Banker, quoting Fuerst, adds the mechanism at the engineering level [11].

This paper’s framework explains what all five describe. AI raises the value of the SaaS bundle because the deployer now sees, in concrete contract terms and operational cost, what the bundle was worth.

The raising of value is not specific to incumbent SaaS. The bundle accrues to any vendor that absorbs the five asymmetries competently, whether a mature SaaS vendor that integrates AI well, an AI-

native SaaS entrant that builds the bundle on new foundations, or a specialist managed-hosting provider that wraps open-weight models with accountability structures.

There are deployment contexts in which an in-house AI build is the better decision, and the paper does not claim otherwise. Greenfield workflows where no SaaS equivalent exists, lightweight internal tools with low compliance exposure, augmentation layers that sit on top of existing SaaS rather than replacing it, and deployments inside sovereignty-mandated frameworks all weaken or eliminate the asymmetries the paper has described. The argument is narrower than “do not replace SaaS with AI”; it is that the asymmetries should be counted before the replacement decision is made in contexts where mature SaaS already exists and is in production.

For the 2026 to 2030 procurement horizon, enterprises replacing mature SaaS with in-house AI builds absorb five asymmetries the SaaSocalypse narrative does not count. This is one risk category among many, and it is one the current narrative systematically misses.

References

- [1] Forrester, “SaaS As We Know It Is Dead: How To Survive The SaaS-pocalypse!” February 2026. <https://www.forrester.com/blogs/saas-as-we-know-it-is-dead-how-to-survive-the-saas-pocalypse/>
- [2] Don Muir, “\$300 Billion Evaporated. The SaaS-Pocalypse Has Begun.” *Forbes*, February 4, 2026. <https://www.forbes.com/sites/donmuir/2026/02/04/300-billion-evaporated-the-saaspocalypse-has-begun/>
- [3] Alex Immerman and Santiago Rodriguez, “Good news: AI Will Eat Application Software.” Andreessen Horowitz, March 2, 2026. <https://a16z.com/good-news-ai-will-eat-application-software/>
- [4] Julien Bek, “Services: The New Software.” Sequoia Capital, March 5, 2026. <https://sequoiacap.com/article/services-the-new-software/>
- [5] Oliver E. Williamson, “Transaction-Cost Economics: The Governance of Contractual Relations.” *Journal of Law and Economics* 22, no. 2 (October 1979): 233–261.
- [6] TermScout, *The 2026 Contract Trust Report*. April 2026. <https://www.term Scout.com/2026-contract-trust-report>
- [7] Regulation (EU) 2024/1689 of the European Parliament and of the Council on artificial intelligence (Artificial Intelligence Act), Article 26 — Obligations of Deployers of High-Risk AI Systems, 2024. <https://artificialintelligenceact.eu/article/26/>
- [8] Cloud Security Alliance, *AI Organizational Responsibilities: Core Security Responsibilities*. April 2024. <https://cloudsecurityalliance.org/research/working-groups/ai-organizational-responsibilities>
- [9] Laura Bratton, “Uber CTO Shows How Claude Code Can Blow Up AI Budgets.” *The Information*, April 14, 2026. <https://www.theinformation.com/newsletters/applied-ai/uber-cto-shows-claude-code-can-blow-ai-budgets>
- [10] Aditya Challapally, Chris Pease, Ramesh Raskar, and Pradyumna Chari, *The GenAI Divide: State of AI in Business 2025*. MIT NANDA, July 2025. https://mlq.ai/media/quarterly_decks/v0.1.State_of_AI_in_Business_2025_Report.pdf
- [11] Steve Banker, “Rethinking the SaaS Apocalypse.” *Forbes*, March 26, 2026. Quoting Chuck Fuerst, “The Seduction of Speed.” <https://www.forbes.com/sites/stevebanker/2026/03/26/rethinking-the-saas-apocalypse/>
- [12] Joy H. Shen and Alex Tamkin, “How AI Impacts Skill Formation.” arXiv preprint arXiv:2601.20245, 2026. <https://arxiv.org/abs/2601.20245>
- [13] Addy Osmani, “Comprehension Debt: The Hidden Cost of AI-Generated Code.” Personal blog, March 2026 (cross-posted at O’Reilly Radar). <https://addyosmani.com/blog/comprehension-debt/>

- [14] Gartner, "Lack of AI-Ready Data Puts AI Projects at Risk." Q&A with Roxane Edjlali, Senior Director Analyst. Gartner Newsroom, Stamford, CT, February 26, 2025. <https://www.gartner.com/en/newsroom/press-releases/2025-02-26-lack-of-ai-ready-data-puts-ai-projects-at-risk>
- [15] MuleSoft, *2025 Connectivity Benchmark Report* (in collaboration with Vanson Bourne and Deloitte Digital), based on interviews with 1,050 IT leaders globally. Salesforce/MuleSoft, 2025. <https://www.mulesoft.com/lp/reports/connectivity-benchmark>
- [16] David Cahn, "AI in 2026: A Tale of Two AIs." Sequoia Capital, December 3, 2025. <https://www.sequoiacap.com/article/ai-in-2026-the-tale-of-two-ais/>